



# Bayesian selection of best subsets via hybrid search

Shiqiang Jin<sup>1</sup> · Gyuhyeong Goh<sup>1</sup>

Received: 22 October 2019 / Accepted: 29 April 2020 / Published online: 11 May 2020  
© Springer-Verlag GmbH Germany, part of Springer Nature 2020

## Abstract

Over the past decades, variable selection for high-dimensional data has drawn increasing attention. With a large number of predictors, there rises a big challenge for model fitting and prediction. In this paper, we develop a new Bayesian method of best subset selection using a hybrid search algorithm that combines a deterministic local search and a stochastic global search. To reduce the computational cost of evaluating multiple candidate subsets for each update, we propose a novel strategy that enables us to calculate exact marginal likelihoods of all neighbor models simultaneously in a single computation. In addition, we establish model selection consistency for the proposed method in the high-dimensional setting in which the number of possible predictors can increase faster than the sample size. Simulation study and real data analysis are conducted to investigate the performance of the proposed method.

**Keywords** Bayesian variable selection · Best subset selection · High-dimensional regression analysis · Hybrid search algorithm

## 1 Introduction

Variable selection plays a key role in recent regression analysis. In many statistical applications, especially in genetics studies, researchers often face situations where the number of candidate predictors is extremely large but the sample size is relatively small, often referred to as high-dimensional regression problems. The most pressing challenge in high-dimensional regression is to identify relevant predictor variables from a large pool of candidates. In an attempt to perform high-dimensional variable selection, a lot of effort has been put into the development of penalized likelihood methods (e.g., Tibshirani 1996; Fan and Li 2001; Zou and Hastie 2005; Zhang 2010). By adding a penalty function to the likelihood criterion, the penalized likelihood

---

✉ Gyuhyeong Goh  
ggoh@ksu.edu

<sup>1</sup> Department of Statistics, Kansas State University, Manhattan, KS 66506, USA

method produces sparse solutions that automatically eliminate the irrelevant predictors from the regression model using the zero-estimates.

In this paper, we are interested in selecting the  $k$  most important predictors out of  $p$  candidates, called best subset selection (Hocking and Leslie 1967). It is well known that best subset selection involves nonconvex optimization problems that are computationally intractable in high-dimensional settings. Although some penalized likelihood approaches such as Least Absolute Shrinkage and Selection Operator (LASSO), Elastic Net (ENET), and Minimax Concave Penalty (MCP) provide a convex surrogate for nonconvex optimization, their applicability to best subset selection is still limited (Bertsimas et al. 2016). Meanwhile, a Bayesian approach to best subset selection, called Bayesian subset regression (BSR), has been proposed by Liang et al. (2013). Using an adaptive Markov chain Monte Carlo (MCMC) algorithm, called the stochastic approximation Monte Carlo (Liang et al. 2007), BSR finds the best subset by performing a stochastic search over the entire model space. However, the stochastic search with a large number of candidate predictors often raises computational challenges including extremely heavy computation and slow convergence. To overcome this limitation, we introduce a hybrid best subset search algorithm that combines a deterministic local search and a stochastic global search in a Bayesian framework. The main attractive feature of our proposed method is that evaluating all possible candidate models for the next update, which is the most expensive part of computation, is simultaneously accomplished in a single computation.

## 2 Basic setup

Consider a multiple linear regression model,

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}, \quad (1)$$

where  $\mathbf{y} = (y_1, \dots, y_n)^\top$  is the  $n$ -dimensional response vector,  $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_p)$  is the  $n \times p$  design matrix with  $\mathbf{x}_j = (x_{1j}, \dots, x_{nj})^\top$ ,  $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)^\top$  is a  $p$ -dimensional coefficient vector, and  $\boldsymbol{\epsilon} = (\epsilon_1, \dots, \epsilon_n)^\top \sim \text{Normal}(\mathbf{0}_n, \sigma^2 \mathbf{I}_n)$ . Throughout this paper, we restrict our attention to high-dimensional regression problems and thus it is always assumed that  $p > n$  and  $\boldsymbol{\beta}$  contains many zero elements. We further assume that  $\mathbf{y}$  and columns of  $\mathbf{X}$  are standardized so that the intercept is excluded from our regression analysis.

In this paper, our goal is to identify the  $k$  most important predictors in (1), where the best subset size,  $k$ , can be considered as being either fixed or varying. To formulate a Bayesian framework for best subset selection, let  $\boldsymbol{\gamma} = \{j : \beta_j \neq 0\}$  be an index set of the active predictors. Given  $\boldsymbol{\gamma}$ , the full model (1) reduces to

$$\mathbf{y} = \mathbf{X}_{\boldsymbol{\gamma}}\boldsymbol{\beta}_{\boldsymbol{\gamma}} + \boldsymbol{\epsilon},$$

where  $\mathbf{X}_{\boldsymbol{\gamma}}$  and  $\boldsymbol{\beta}_{\boldsymbol{\gamma}}$  are a sub-matrix of  $\mathbf{X}$  and a sub-vector of  $\boldsymbol{\beta}$  that are determined by  $\boldsymbol{\gamma}$ , respectively. For algebraic and computational convenience, given  $\boldsymbol{\gamma}$ , we consider conjugate priors for  $\boldsymbol{\beta}_{\boldsymbol{\gamma}}$  and  $\sigma^2$  as follows:

$$\begin{aligned}\beta_{\mathcal{Y}}|\sigma^2, \mathcal{Y} &\sim \text{Normal}(0, \tau\sigma^2\mathbf{I}_{|\mathcal{Y}|}), \\ \sigma^2 &\sim \text{Inverse-Gamma}(a_{\sigma}/2, b_{\sigma}/2),\end{aligned}$$

where  $\tau$ ,  $a_{\sigma}$ , and  $b_{\sigma}$  are hyperparameters and  $|\mathcal{Y}|$  denotes the number of elements in the set  $\mathcal{Y}$ . To impose the constraint  $|\mathcal{Y}| = k$ , we define the prior distribution of  $\mathcal{Y}$  by  $\pi(\mathcal{Y}) \propto \mathbb{I}(|\mathcal{Y}| = k)$ , where  $\mathbb{I}(\cdot)$  is an indicator function. Let  $m(\mathbf{y}|\mathcal{Y})$  be the marginal likelihood given  $\mathcal{Y}$ . Using the kernels of normal density and inverse gamma density, the marginal likelihood can be calculated as

$$\begin{aligned}m(\mathbf{y}|\mathcal{Y}) &= \int f(\mathbf{y}|\beta_{\mathcal{Y}}, \sigma^2)\pi(\beta_{\mathcal{Y}}|\sigma^2, \mathcal{Y})\pi(\sigma^2)d\beta_{\mathcal{Y}}d\sigma^2 \\ &\propto \frac{(\tau^{-1})^{\frac{|\mathcal{Y}|}{2}}}{|\mathbf{X}_{\mathcal{Y}}^{\top}\mathbf{X}_{\mathcal{Y}} + \tau^{-1}\mathbf{I}_{|\mathcal{Y}|}|^{\frac{1}{2}}(\mathbf{y}^{\top}\mathbf{H}_{\mathcal{Y}}\mathbf{y} + b_{\sigma})^{\frac{a_{\sigma}+n}{2}}},\end{aligned}\quad (2)$$

where  $f(\mathbf{y}|\beta_{\mathcal{Y}}, \sigma^2)$  denotes the reduced model likelihood function given  $\mathcal{Y}$  and  $\mathbf{H}_{\mathcal{Y}} = \mathbf{I}_n - \mathbf{X}_{\mathcal{Y}}(\mathbf{X}_{\mathcal{Y}}^{\top}\mathbf{X}_{\mathcal{Y}} + \tau^{-1}\mathbf{I}_{|\mathcal{Y}|})^{-1}\mathbf{X}_{\mathcal{Y}}^{\top}$ .

From Bayes' theorem, the posterior model probability of  $\mathcal{Y}$  is proportional to  $\pi(\mathcal{Y}|\mathbf{y}) \propto m(\mathbf{y}|\mathcal{Y})\pi(\mathcal{Y}) \propto m(\mathbf{y}|\mathcal{Y})\mathbb{I}(|\mathcal{Y}| = k)$ . Therefore, Bayesian best subset selection can be performed by maximizing  $m(\mathbf{y}|\mathcal{Y})$  over  $\mathcal{Y}$  subject to the constraint  $|\mathcal{Y}| = k$ . Keep in mind that high-dimensional and non-convex optimization problems arise in our framework.

**Remark 1** After selecting the best model  $\mathcal{Y}$ , Bayesian inferences about  $\beta_{\mathcal{Y}}$  and  $\sigma^2$  can be easily made by drawing samples from  $\pi(\beta_{\mathcal{Y}}, \sigma^2|\mathbf{y}, \mathcal{Y})$ . Due to the conjugacy of the priors for  $\beta_{\mathcal{Y}}$  and  $\sigma^2$ , the posterior sampling can be implemented via Gibbs sampling with the closed-form full-conditional distributions.

### 3 Best subset selection with a fixed-size

In this section, we develop a new Bayesian approach to best subset selection with a fixed subset size  $k$ . Let  $\hat{\mathcal{Y}}$  be the current estimate of the best subset of size  $k$ . Our strategy is to update  $\hat{\mathcal{Y}}$  by searching over neighbor models iteratively. To this end, for an index set  $\mathcal{Y}$ , define  $\mathcal{N}_{+}(\mathcal{Y}) = \{\mathcal{Y} \cup \{i\} : i \notin \mathcal{Y}\}$ , which represents the set of larger neighbors of  $\mathcal{Y}$  obtained by adding a new index to  $\mathcal{Y}$ . Similarly, define  $\mathcal{N}_{-}(\mathcal{Y}) = \{\mathcal{Y} \setminus \{j\} : j \in \mathcal{Y}\}$  to be the set of smaller neighbors of  $\mathcal{Y}$  obtained by deleting an index from  $\mathcal{Y}$ .

We introduce a deterministic search algorithm in Algorithm 1 that converges to a local maximum of  $m(\mathbf{y}|\mathcal{Y})$  subject to  $|\mathcal{Y}| = k$ . The following theorem proves the convergence of the proposed deterministic search algorithm.

**Theorem 1** *The deterministic search in Algorithm 1 monotonically increases the objective function,  $m(\mathbf{y}|\mathcal{Y})$ , subject to  $|\mathcal{Y}| = k$ . In addition, the algorithm terminates in a finite number of iterations.*

**Algorithm 1** Deterministic best subset search with a fixed  $k$ 

1. Define  $\hat{\mathbf{y}}$  to be an initial subset of size  $k$ ; and set  $\hat{\mathbf{y}}^{(0)} = \hat{\mathbf{y}}$ .
2. **Repeat** for  $t = 1, 2, \dots$ ,
  - a) Compute  $\tilde{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y})$ ;  
# i.e., select the locally best subset of size  $k+1$
  - b) Compute  $\hat{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})} m(\mathbf{y}|\mathbf{y})$ ;  
# i.e., select the locally best subset of size  $k$
  - c) Update  $\hat{\mathbf{y}} \leftarrow \hat{\mathbf{y}}^{(t)}$ ;
- until**  $\hat{\mathbf{y}}^{(t-1)} = \hat{\mathbf{y}}^{(t)}$ . # i.e., terminate immediately if no update is made
3. Return  $\hat{\mathbf{y}}$ .

**Algorithm 2** Stochastic best subset search with a fixed  $k$ 

1. Define  $\hat{\mathbf{y}}$  to be an initial subset of size  $k$ ; and set  $\hat{\mathbf{y}}^{(0)} = \hat{\mathbf{y}}$ .
2. **Repeat** for  $t = 1, \dots, T$ :
  - (a) Update  $\tilde{\mathbf{y}}^{(t)}$  by selecting one from  $\mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})$  with probability

$$\frac{m(\mathbf{y}|\mathbf{y})}{\sum_{\mathbf{y}' \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y}')} , \quad \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)});$$

- (b) Update  $\hat{\mathbf{y}}^{(t)}$  by selecting one from  $\mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})$  with probability

$$\frac{m(\mathbf{y}|\mathbf{y})}{\sum_{\mathbf{y}' \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})} m(\mathbf{y}|\mathbf{y}')} , \quad \mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)});$$

- (c) **If**  $m(\mathbf{y}|\hat{\mathbf{y}}) < m(\mathbf{y}|\hat{\mathbf{y}}^{(t)})$ , **then** update  $\hat{\mathbf{y}} \leftarrow \hat{\mathbf{y}}^{(t)}$ .

3. Return  $\hat{\mathbf{y}}$ .

**Proof of Theorem 1** Let  $\hat{\mathbf{y}}^{(t-1)}$  be the best subset of size  $k$  updated by the  $t - 1$ -th iteration. Then,  $\hat{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})} m(\mathbf{y}|\mathbf{y})$ , where

$$\tilde{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y}).$$

Since  $\hat{\mathbf{y}}^{(t-1)}$  also belongs to  $\mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})$ , we thus have  $m(\mathbf{y}|\hat{\mathbf{y}}^{(t-1)}) \leq m(\mathbf{y}|\hat{\mathbf{y}}^{(t)})$ , which proves our first statement. Since the number of all possible states of  $\mathbf{y}$  satisfying  $|\mathbf{y}| = k$  is finite, the algorithm terminates in a finite number of iterations. This completes our proof.  $\square$

Although the deterministic search algorithm converges rapidly, a possible drawback is that the algorithm can get trapped in a local optimum. As an alternative, we introduce a stochastic search algorithm in Algorithm 2 that converges to a global maximum of  $m(\mathbf{y}|\mathbf{y})$  with the constraint  $|\mathbf{y}| = k$ . Note that the proposed stochastic search algorithm generates a Markov chain,  $\{\hat{\mathbf{y}}^{(t)}, t = 1, \dots, T\}$ , with the state space  $\{\mathbf{y} : |\mathbf{y}| = k\}$ . Hence, if the current estimate  $\hat{\mathbf{y}}$  is not the global maximum, then we must observe

**Algorithm 3** Hybrid best subset search with a fixed  $k$ 

1. Define an initial model of size  $k$  as  $\hat{\mathbf{y}}$ .
2. Set  $\hat{\mathbf{y}}^{(0)} = \hat{\mathbf{y}}$ .
3. **Repeat** for  $t = 1, 2, \dots$ : # Deterministic Search Step
  - (a) Compute  $\tilde{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y})$ ;
  - (b) Compute  $\hat{\mathbf{y}}^{(t)} = \arg \max_{\mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})} m(\mathbf{y}|\mathbf{y})$ ;
  - (c) Update  $\hat{\mathbf{y}} \leftarrow \hat{\mathbf{y}}^{(t)}$ ;
- until**  $\hat{\mathbf{y}}^{(t-1)} = \hat{\mathbf{y}}^{(t)}$ .
4. Set  $\hat{\mathbf{y}}^{(0)} = \hat{\mathbf{y}}$  and choose a tuning parameter  $\alpha \in (0, 1]$ .
5. **Repeat** for  $t = 1, \dots, T$ : # Stochastic Search Step
  - (a) Update  $\tilde{\mathbf{y}}^{(t)}$  by drawing one from  $\mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})$  with probability

$$\frac{m(\mathbf{y}|\mathbf{y})^\alpha}{\sum_{\mathbf{y}' \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y}')^\alpha}, \quad \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)}); \quad (3)$$

- (b) Update  $\hat{\mathbf{y}}^{(t)}$  by drawing one from  $\mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})$  with probability

$$\frac{m(\mathbf{y}|\mathbf{y})^\alpha}{\sum_{\mathbf{y}' \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)})} m(\mathbf{y}|\mathbf{y}')^\alpha}, \quad \mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}}^{(t)}); \quad (4)$$

- (c) **If**  $m(\mathbf{y}|\hat{\mathbf{y}}) < m(\mathbf{y}|\hat{\mathbf{y}}^{(t)})$ , **then** update  $\hat{\mathbf{y}} \leftarrow \hat{\mathbf{y}}^{(t)}$  and immediately go to Step 2.

6. Return  $\hat{\mathbf{y}}$ .

that  $m(\mathbf{y}|\hat{\mathbf{y}}) < m(\mathbf{y}|\hat{\mathbf{y}}^{(t)})$  with probability one as  $T \rightarrow \infty$ . Therefore, as we iterate the algorithm,  $\hat{\mathbf{y}}$  converges to the global maximum.

Although the stochastic search algorithm eventually reaches the global optimum, it requires the large number of iterations, which is computationally inefficient. To develop a computationally efficient global optimization algorithm, we propose a hybrid best subset selection algorithm that combines the stochastic global search algorithm and the deterministic local search algorithm into one. The proposed hybrid search algorithm is given in Algorithm 3. In the proposed hybrid search algorithm, the deterministic search is first used to find a local optimum in an efficient manner. Then, the stochastic search is employed to check whether or not the deterministic search algorithm has reached the global optimum. Note that if the stochastic search finds a better subset  $\hat{\mathbf{y}}^{(t)}$  than the current best subset  $\hat{\mathbf{y}}$  (i.e., when  $m(\mathbf{y}|\hat{\mathbf{y}}) < m(\mathbf{y}|\hat{\mathbf{y}}^{(t)})$  occurs), then we immediately stop the stochastic search procedure and go back to the deterministic search step with the updated state of  $\hat{\mathbf{y}}$ . To improve the performance of stochastic search, we introduce a tuning parameter  $\alpha \in (0, 1]$  which acts as a precision parameter. For example, in Step 5a of Algorithm 3, we select one among  $\{\mathbf{y} : \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})\}$  with probabilities  $\{m(\mathbf{y}|\mathbf{y})^\alpha / \sum_{\mathbf{y}' \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y}')^\alpha : \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})\}$ , whereas, in Step 2a of Algorithm 2, we select one among  $\{\mathbf{y} : \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})\}$  with probabilities  $\{m(\mathbf{y}|\mathbf{y}) / \sum_{\mathbf{y}' \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})} m(\mathbf{y}|\mathbf{y}') : \mathbf{y} \in \mathcal{N}_+(\hat{\mathbf{y}}^{(t-1)})\}$ . As  $\alpha \rightarrow 0$ , the probability distributions in (3) and (4) will be more spread out so that the chance of getting stuck in the local maximum will be reduced in the stochastic search step. If we set  $\alpha = 1$ ,

the stochastic search step is analogous to Algorithm 2. In our simulation study and real data analysis, we set  $\alpha = \min\{1, \log(2)/\log(m_{(1)}/m_{(2)})\}$ , where  $m_{(1)}$  and  $m_{(2)}$  are, respectively, the first and second largest values of  $\{m(\mathbf{y}|\boldsymbol{\gamma}) : \boldsymbol{\gamma} \in \mathcal{N}_+(\hat{\boldsymbol{\gamma}})\}$ .

Note that  $m(\mathbf{y}|\hat{\boldsymbol{\gamma}}) < m(\mathbf{y}|\hat{\boldsymbol{\gamma}}^{(t)})$  if and only if  $m(\mathbf{y}|\hat{\boldsymbol{\gamma}})^\alpha < m(\mathbf{y}|\hat{\boldsymbol{\gamma}}^{(t)})^\alpha$  for any  $\alpha \in (0, 1]$ . Hence, as in Algorithm 2, if the current estimate  $\hat{\boldsymbol{\gamma}}$  is not the global maximum, then in Step 5 of Algorithm 3 we must observe that  $m(\mathbf{y}|\hat{\boldsymbol{\gamma}}) < m(\mathbf{y}|\hat{\boldsymbol{\gamma}}^{(t)})$  with probability one as  $T \rightarrow \infty$ . Therefore, we can reach the global maximum by repeating Algorithm 3 with sufficiently large  $T$ .

In general, calculating marginal likelihoods for many candidate models, which is a necessary step in Bayesian model selection, is computationally expensive, even if an explicit form is available as in (2). The great merit of the proposed method is that evaluating the marginal likelihoods of all the candidates for the next update can be done simultaneously *within a single computation*. To this end, let  $\hat{\boldsymbol{\gamma}}$  be an index set of a model of size  $k$  and  $\tilde{\boldsymbol{\gamma}}$  be an index set of a model of size  $k + 1$ . For any  $i \notin \hat{\boldsymbol{\gamma}}$ , it can be shown that

$$m(\mathbf{y}|\hat{\boldsymbol{\gamma}} \cup \{i\}) \propto \left\{ \mathbf{y}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y} - \frac{(\mathbf{x}_i^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y})^2}{\tau^{-1} + \mathbf{x}_i^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{x}_i} + b_\sigma \right\}^{-\frac{a_\sigma + n}{2}} \times (\tau^{-1} + \mathbf{x}_i^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{x}_i)^{-1/2}, \quad (5)$$

where  $\mathbf{x}_i$  is the  $i$ -th column of  $\mathbf{X}$ . The proof of (5) is given in Appendix A. Note that  $\mathbf{x}_i^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{x}_i$  is the  $i$ -th diagonal element of  $\mathbf{X}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{X}$  and that  $\mathbf{x}_i^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y}$  is the  $i$ -th element of  $\mathbf{X}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y}$ . This implies that Eq. (5) is the  $i$ -th element of the following  $p$ -dimensional vector:

$$\mathbf{m}_+(\hat{\boldsymbol{\gamma}}) = \left\{ (\mathbf{y}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y} + b_\sigma) \mathbf{1}_p - \frac{(\mathbf{X}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{y})^2}{\frac{1}{\tau} \mathbf{1}_p + \text{diag}(\mathbf{X}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{X})} \right\}^{-\frac{a_\sigma + n}{2}} \times \left\{ \frac{1}{\tau} \mathbf{1}_p + \text{diag}(\mathbf{X}^\top \mathbf{H}_{\hat{\boldsymbol{\gamma}}} \mathbf{X}) \right\}^{-1/2}, \quad (6)$$

where  $\mathbf{a}^x$  and  $\mathbf{a}/\mathbf{b}$  indicate entrywise operations for generic vectors  $\mathbf{a}$  and  $\mathbf{b}$ , accordingly. For example,  $\mathbf{a}^x = (a_1^x, \dots, a_p^x)$  and  $\mathbf{a}/\mathbf{b} = (a_1/b_1, \dots, a_p/b_p)$ . Note that  $\mathcal{N}_+(\hat{\boldsymbol{\gamma}}) = \{\hat{\boldsymbol{\gamma}} \cup \{i\} : i \notin \hat{\boldsymbol{\gamma}}\}$ . Therefore, evaluating  $m(\mathbf{y}|\boldsymbol{\gamma})$  for all  $\boldsymbol{\gamma} \in \mathcal{N}_+(\hat{\boldsymbol{\gamma}})$  can be done simultaneously in a single computation by obtaining the sub-vector of  $\mathbf{m}_+(\hat{\boldsymbol{\gamma}})$  for  $\{i : i \notin \hat{\boldsymbol{\gamma}}\}$ . Similarly, for any  $j \in \tilde{\boldsymbol{\gamma}}$ , we can show that

$$m(\mathbf{y}|\tilde{\boldsymbol{\gamma}} \setminus \{j\}) \propto \left\{ \mathbf{y}^\top \mathbf{H}_{\tilde{\boldsymbol{\gamma}}} \mathbf{y} + \frac{(\mathbf{x}_j^\top \mathbf{H}_{\tilde{\boldsymbol{\gamma}}} \mathbf{y})^2}{\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\boldsymbol{\gamma}}} \mathbf{x}_j} + b_\sigma \right\}^{-\frac{a_\sigma + n}{2}} \times (\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\boldsymbol{\gamma}}} \mathbf{x}_j)^{-1/2}. \quad (7)$$

The proof of (7) is given in Appendix B. Define

$$\mathbf{m}_-(\tilde{\mathbf{y}}) = \left\{ (\mathbf{y}^\top \mathbf{H}_{\tilde{\mathbf{y}}} \mathbf{y} + b_\sigma) \mathbf{1}_p + \frac{(\mathbf{X}^\top \mathbf{H}_{\tilde{\mathbf{y}}} \mathbf{y})^2}{\frac{1}{\tau} \mathbf{1}_p - \text{diag}(\mathbf{X}^\top \mathbf{H}_{\tilde{\mathbf{y}}} \mathbf{X})} \right\}^{-\frac{a\sigma + n}{2}} \times \left\{ \frac{1}{\tau} \mathbf{1}_p - \text{diag}(\mathbf{X}^\top \mathbf{H}_{\tilde{\mathbf{y}}} \mathbf{X}) \right\}^{-1/2}. \quad (8)$$

It is easy to check that Eq. (7) is the  $j$ -th element of  $\mathbf{m}_-(\tilde{\mathbf{y}})$ . Since  $\mathcal{N}_-(\tilde{\mathbf{y}}) = \{\tilde{\mathbf{y}} \setminus \{j\} : j \in \tilde{\mathbf{y}}\}$ , evaluating  $m(\mathbf{y}|\mathbf{y})$  for all  $\mathbf{y} \in \mathcal{N}_-(\tilde{\mathbf{y}})$  can be done simultaneously in a single computation by obtaining the sub-vector of  $\mathbf{m}_-(\tilde{\mathbf{y}})$  for  $\{j : j \in \tilde{\mathbf{y}}\}$ . Using the aforementioned calculation strategy, we can easily and quickly implement steps 3a), 3b), 5a), and 5b) of Algorithm 3. It is also worth noting that the proposed calculation method enables us to avoid multiple computations of inverse and determinant of matrices in Eq. (2).

## 4 Best subset selection with varying sizes

In this section, we extend the proposed method to handle the case of varying  $k \leq K$ , where  $K$  is a prespecified upper bound. This is a common setting for high-dimensional best subset selection problems (e.g., Bertsimas et al. 2016; Liang et al. 2013). In a Bayesian framework, this extension can be easily done by assigning an appropriate prior for unknown  $k$ . As a non-informative prior, one may consider a discrete uniform prior for  $k$ , that is,  $k \sim \text{Uniform}\{1, \dots, K\}$ . However, the uniform prior tends to assign larger probability to a larger subset due to the fact that the total number of subsets of size  $k$  is  $\binom{p}{k}$ , which tends to increase exponentially as  $k$  increases. To resolve this issue, using a similar idea of Chen and Chen (2008), we consider

$$\pi(k) \propto \mathbb{I}(k \leq K) / \binom{p}{k}.$$

From Bayes' theorem, Bayesian best subset selection is then performed by maximizing

$$\pi(\mathbf{y}|\mathbf{y}) \propto m(\mathbf{y}|\mathbf{y}) \mathbb{I}(|\mathbf{y}| \leq K) / \binom{p}{|\mathbf{y}|}. \quad (9)$$

Note that (9) is equivalent to the following optimization problem:

$$\max_{\mathbf{y}} \left\{ m(\mathbf{y}|\mathbf{y}) \mathbb{I}(|\mathbf{y}| = k) / \binom{p}{k} \right\} \quad \text{subject to } k \leq K.$$

Therefore, the optimization problem in Eq. (9) can be solved by proceeding the following steps:

1. For  $k = 1, \dots, K$ , compute

$$\hat{\boldsymbol{y}}_k = \arg \max_{\boldsymbol{y}: |\boldsymbol{y}|=k} m(\boldsymbol{y}|\boldsymbol{y})$$

using Algorithm 3.

2. Compute

$$\hat{k} = \arg \max_{1 \leq k \leq K} \left\{ \log m(\boldsymbol{y}|\hat{\boldsymbol{y}}_k) - \log \binom{p}{k} \right\}.$$

3. Return  $\hat{\boldsymbol{y}} = \hat{\boldsymbol{y}}_{\hat{k}}$ .

Note that if a parallel or cluster computing environment is available, then it can be applied to the first step of the above procedure to run the for-loop over  $k$  in parallel. The following theorem shows that the proposed Bayesian approach achieves the model selection consistency in the high-dimensional setting where  $p > n$ .

**Theorem 2** Define  $\Gamma = \{\boldsymbol{y} : |\boldsymbol{y}| \leq K, \boldsymbol{y} \neq \boldsymbol{y}_*\}$ , where  $\boldsymbol{y}_*$  is the true model. Assume that  $p = O(n^\xi)$  for  $\xi \geq 1$ . Under the asymptotic identifiability condition of Chen and Chen (2008), if  $\tau \rightarrow \infty$  as  $n \rightarrow \infty$  but  $\tau = o(n)$ , then the proposed Bayesian subset selection possesses the Bayesian model selection consistency, that is,

$$\pi(\boldsymbol{y}_*|\boldsymbol{y}) > \max_{\boldsymbol{y} \in \Gamma} \pi(\boldsymbol{y}|\boldsymbol{y}) \quad (10)$$

in probability as  $n \rightarrow \infty$ .

**Proof of Theorem 2** From the Laplace approximation of Kass and Raftery (1995), we have

$$\log m(\boldsymbol{y}|\boldsymbol{y}) = \log f(\boldsymbol{y}|\hat{\boldsymbol{\beta}}_{\boldsymbol{y}}, \hat{\sigma}^2) - \frac{|\boldsymbol{y}|}{2} \log n + o_p(\log n),$$

where  $\hat{\boldsymbol{\beta}}_{\boldsymbol{y}} = (\mathbf{X}_{\boldsymbol{y}}^\top \mathbf{X}_{\boldsymbol{y}})^{-1} \mathbf{X}_{\boldsymbol{y}}^\top \boldsymbol{y}$  and  $\hat{\sigma}^2 = \|\mathbf{H}_{\boldsymbol{y}} \boldsymbol{y}\|^2/n$ . Ignoring a smaller order term than  $\log n$ , our posterior criterion (9) can be approximated as

$$\begin{aligned} -2 \log \pi(\boldsymbol{y}|\boldsymbol{y}) &\approx -2 \log f(\boldsymbol{y}|\hat{\boldsymbol{\beta}}_{\boldsymbol{y}}, \hat{\sigma}^2) \\ &\quad + |\boldsymbol{y}| \log n + 2 \log \binom{p}{k}, \end{aligned} \quad (11)$$

which is equivalent to the extended Bayesian Information Criterion (BIC) (Chen and Chen 2008). Therefore, it follows from Theorem 1 of Chen and Chen (2008) that Eq. (10) holds in probability.  $\square$



## 5 Simulation study

### 5.1 Model selection performance comparison

In this section, we investigate the variable selection performance of our best subset selection algorithm on simulated high-dimensional data. We consider three scenarios of data generating processes.

In Scenario I, we generate the data  $\{(y_i, x_{i1}, \dots, x_{ip}) : i = 1, \dots, n\}$  from the following linear regression model:

$$y_i = \sum_{j=1}^p \beta_j x_{ij} + \epsilon_i \quad (12)$$

where  $n = 100$ ,  $(x_{i1}, \dots, x_{ip})^\top \stackrel{iid}{\sim} \text{Normal}(\mathbf{0}_p, \mathbf{\Sigma})$  with  $\mathbf{\Sigma} = (\Sigma_{ij})_{p \times p}$  and  $\Sigma_{ij} = \rho^{|i-j|}$ ,  $(\epsilon_1, \dots, \epsilon_n) \sim \text{Normal}(0, \mathbf{I}_n)$ , four  $\beta_j$ 's are randomly selected and then generated from  $\text{Uniform}\{-1, -2, 1, 2\}$  independently, and the remaining  $\beta$ -coefficients are set equal to 0.

In Scenario II, we consider the same setting as in Scenario I, but we use the following correlated error structure instead of the independent errors:  $(\epsilon_1, \dots, \epsilon_n)^\top \sim \text{Normal}(0, \mathbf{V})$  with  $\mathbf{V} = (v_{ij})_{n \times n}$  and  $v_{ij} = 0.5^{|i-j|}$ .

In Scenario III, we consider the same setting as in Scenario I, but the heteroscedastic error structure and the inverse-wishart random variance for predictors are used as follows:  $\epsilon_i \stackrel{ind}{\sim} \text{Normal}(0, \sigma_i^2)$  with  $\sigma_i^2 = 1 + 0.5\sqrt{x_{i1}}$  and  $\mathbf{\Sigma} \sim \text{Inverse-Wishart}((1000 - 1)\mathbf{\Psi}, p + 1000)$  with  $\mathbf{\Psi} = (\Psi_{ij})_{p \times p}$  and  $\Psi_{ij} = \rho^{|i-j|}$ .

For each scenario, we consider four different cases: (i)  $p = 200$ ,  $\rho = 0.1$ , (ii)  $p = 200$ ,  $\rho = 0.9$ , (iii)  $p = 1000$ ,  $\rho = 0.1$ , and (iv)  $p = 1000$ ,  $\rho = 0.9$ . We assume that there is no prior information. Hence, to make our prior distributions non-informative, we set  $a_\sigma = b_\sigma = 1$  and  $\tau = (\log p)^2$ , which satisfies the sufficient condition for model selection consistency discussed in Theorem 2. We further assume that the true number of active predictors is unknown. Hence, we employ the proposed method in Section 4 for varying  $k$  with  $K = \lceil n^{2/3} \rceil = 22$ . For the hybrid search algorithm given in Algorithm 3, we set  $T = 100$ , which is a remarkably small number of iterations compared to existing stochastic search algorithms (e.g., Kirkpatrick et al. 1983; George and McCulloch 1993; Hans et al. 2007), and marginal correlations between the response and each predictor are used to define the initial value of  $\hat{\mathbf{y}}$ . For comparison purposes, we also employ four popular penalized likelihood methods: LASSO (Tibshirani 1996), ENET (Zou and Hastie 2005), Smoothly Clipped Absolute Deviation (SCAD) (Fan and Li 2001), and MCP (Zhang 2010), where the regularization parameters or tuning parameters are determined by the extended BIC in (11) to achieve a fair comparison with the proposed method. In this simulation study, LASSO and ENET are implemented by R package `glmnet` and MCP and SCAD are performed by R package `ncvreg`. To evaluate the variable selection performance, we calculate the false discovery rate, the percentage of selecting the exact true model, the average size of the selected model, and the mean of Hamming distance based on 2000

**Table 1** Simulation study results based on 2000 Monte Carlo replications for Scenario I with Cases i–iv

Case	Method	FDR (s.e.)	TRUE (s.e.)	SIZE (s.e.)	HAM (s.e.)
i	Proposed	0.006 (0.001)	96.900 (0.388)	4.032 (0.004)	0.032 (0.004)
	SCAD	0.034 (0.002)	85.200 (0.794)	4.188 (0.011)	0.188 (0.011)
	MCP	0.035 (0.002)	84.750 (0.804)	4.191 (0.011)	0.191 (0.011)
	ENET	0.016 (0.001)	92.700 (0.582)	4.087 (0.007)	0.087 (0.007)
	LASSO	0.020 (0.002)	91.350 (0.629)	4.109 (0.009)	0.109 (0.009)
ii	Proposed	0.023 (0.002)	88.750 (0.707)	3.985 (0.006)	0.203 (0.014)
	SCAD	0.059 (0.003)	74.150 (0.979)	4.107 (0.015)	0.480 (0.022)
	MCP	0.137 (0.004)	55.400 (1.112)	4.264 (0.020)	1.098 (0.034)
	ENET	0.501 (0.004)	0.300 (0.122)	7.716 (0.072)	5.018 (0.052)
	LASSO	0.276 (0.004)	15.550 (0.811)	5.308 (0.033)	2.038 (0.034)
iii	Proposed	0.004 (0.001)	98.100 (0.305)	4.020 (0.003)	0.020 (0.003)
	SCAD	0.027 (0.002)	87.900 (0.729)	4.145 (0.010)	0.145 (0.010)
	MCP	0.031 (0.002)	86.550 (0.763)	4.172 (0.013)	0.172 (0.013)
	ENET	0.035 (0.002)	84.850 (0.802)	4.181 (0.013)	0.206 (0.012)
	LASSO	0.014 (0.001)	93.850 (0.537)	4.073 (0.007)	0.073 (0.007)
iv	Proposed	0.023 (0.002)	89.850 (0.675)	4.005 (0.005)	0.190 (0.013)
	SCAD	0.068 (0.003)	74.250 (0.978)	4.196 (0.014)	0.493 (0.023)
	MCP	0.152 (0.004)	53.750 (1.115)	4.226 (0.017)	1.202 (0.035)
	ENET	0.417 (0.005)	0.150 (0.087)	6.228 (0.068)	4.089 (0.043)
	LASSO	0.265 (0.004)	19.500 (0.886)	5.139 (0.029)	1.909 (0.035)

Notation: FDR—false discovery rate; TRUE—percentage of the true model detected; SIZE—selected model size; HAM—Hamming distance; s.e.— standard error

Monte Carlo replications. False discovery rate is defined as  $|\hat{\mathcal{Y}} \setminus \mathcal{Y}_*|/|\hat{\mathcal{Y}}|$ , and Hamming distance is defined as  $|\hat{\mathcal{Y}} \setminus \mathcal{Y}_*| + |\mathcal{Y}_* \setminus \hat{\mathcal{Y}}|$ , where  $\mathcal{Y}_*$  denotes the index set of the true model.

The results are shown in Tables 1, 2, and 3. In each scenario, the proposed method outperforms all the penalized likelihood methods for all three cases. The proposed method always achieves the smallest false discovery rate and this implies that the proposed method provides the minimum proportion of incorrectly identifying the true active predictors as inactive. According to the selected model size and the Hamming distance, we argue that the proposed method tends to select the closest model to the true model. In addition, the proposed method selects the exact true model with high probability. Hence, this finite sample performance supports our theoretical result in Theorem 2.

## 5.2 Computational speed comparison

In this section, we investigate the computational speed improvement of the proposed method. We are specifically interested in assessing the computational efficiency of the simultaneous marginal likelihood evaluation using the proposed formulas of (6) and

**Table 2** Simulation study results based on 2000 Monte Carlo replications for Scenario II with Cases i–iv

Case	Method	FDR (s.e.)	TRUE (s.e.)	SIZE (s.e.)	HAM (s.e.)
i	Proposed	0.017 (0.001)	92.150 (0.602)	4.091 (0.008)	0.092 (0.008)
	SCAD	0.032 (0.002)	85.850 (0.780)	4.176 (0.011)	0.175 (0.011)
	MCP	0.033 (0.002)	85.500 (0.780)	4.178 (0.011)	0.177 (0.011)
	ENET	0.020 (0.001)	91.250 (0.632)	4.104 (0.008)	0.104 (0.008)
	LASSO	0.019 (0.001)	91.450 (0.625)	4.101 (0.008)	0.101 (0.008)
ii	Proposed	0.030 (0.002)	85.950 (0.777)	4.030 (0.007)	0.234 (0.014)
	SCAD	0.058 (0.003)	75.100 (0.967)	4.088 (0.013)	0.471 (0.022)
	MCP	0.123 (0.004)	58.800 (1.101)	4.218 (0.019)	0.983 (0.032)
	ENET	0.496 (0.004)	0.300 (0.000)	7.636 (0.072)	4.965 (0.051)
	LASSO	0.269 (0.004)	17.600 (0.852)	5.235 (0.033)	1.987 (0.035)
iii	Proposed	0.016 (0.001)	92.550 (0.587)	4.083 (0.007)	0.084 (0.007)
	SCAD	0.028 (0.002)	87.800 (0.732)	4.152 (0.010)	0.152 (0.010)
	MCP	0.032 (0.002)	86.200 (0.771)	4.191 (0.016)	0.191 (0.016)
	ENET	0.030 (0.002)	85.850 (0.780)	4.149 (0.012)	0.183 (0.011)
	LASSO	0.017 (0.001)	92.450 (0.591)	4.088 (0.008)	0.090 (0.008)
iv	Proposed	0.029 (0.002)	88.100 (0.724)	4.043 (0.006)	0.219 (0.015)
	SCAD	0.058 (0.003)	77.650 (0.932)	4.152 (0.012)	0.419 (0.021)
	MCP	0.148 (0.004)	54.700 (1.113)	4.245 (0.019)	1.169 (0.036)
	ENET	0.427 (0.005)	0.200 (0.100)	6.379 (0.069)	4.204 (0.044)
	LASSO	0.263 (0.004)	19.050 (0.878)	5.157 (0.028)	1.868 (0.033)

Notation: FDR—false discovery rate; TRUE—percentage of the true model detected; SIZE—selected model size; HAM—Hamming distance; s.e.— standard error

(8). To this end, we generate the data  $\{(y_i, x_{i1}, \dots, x_{ip}) : i = 1, \dots, n\}$  using the data generating process of Scenario I in Sect. 5.1 for  $n = 100$ ,  $\rho = 0.5$ , and varying  $p$ .


The execution time is measured by implementing Algorithm 3 with  $k = 4$  which is the number of the true active predictors. For the purpose of comparison, we also measure the execution time of Algorithm 3 after replacing the proposed simultaneous marginal likelihood evaluation with the for-loop, which is commonly used in Bayesian computation, for evaluating the marginal likelihoods of candidate models for the next move in Steps 3a, 3b, 5a, and 5b. The simulation is implemented by R using a Linux-based desktop computer with an Intel Core i5-2400 processor and 16 GB of memory.

Figure 1 displays the trend of the average execution time as  $p$  increases. The proposed method has a constant computational speed for increasing values of  $p$ , while the average execution time of the for-loop method steadily increases as the number of candidate predictors grows. For example, when  $p = 100$ , the proposed method takes only 0.122 s until convergence occurs, whereas the for-loop requires 4.020 s to reach the convergence. When  $p = 1000$ , the average execution time of the proposed method is 0.208 s while the for-loop requires 20.289 s on average to complete the task. The simulation result clearly shows that the proposed method is computationally much faster and more efficient than the use of the for-loop.

**Table 3** Simulation study results based on 2000 Monte Carlo replications for Scenario III with Cases i–iv

Case	Method	FDR (s.e.)	TRUE (s.e.)	SIZE (s.e.)	HAM (s.e.)
i	Proposed	0.006 (0.001)	96.850 (0.391)	4.031 (0.004)	0.032 (0.004)
	SCAD	0.033 (0.002)	85.450 (0.789)	4.176 (0.011)	0.175 (0.011)
	MCP	0.034 (0.002)	85.150 (0.795)	4.184 (0.011)	0.184 (0.011)
	ENET	0.028 (0.002)	87.250 (0.746)	4.148 (0.010)	0.156 (0.010)
	LASSO	0.020 (0.002)	90.600 (0.653)	4.102 (0.008)	0.104 (0.008)
ii	Proposed	0.042 (0.002)	80.150 (0.892)	3.952 (0.008)	0.385 (0.019)
	SCAD	0.098 (0.004)	62.300 (1.084)	4.132 (0.017)	0.803 (0.029)
	MCP	0.180 (0.005)	44.050 (1.110)	4.282 (0.023)	1.470 (0.038)
	ENET	0.451 (0.005)	0.350 (0.132)	6.287 (0.071)	4.429 (0.046)
	LASSO	0.312 (0.004)	10.800 (0.694)	5.206 (0.036)	2.416 (0.036)
iii	Proposed	0.004 (0.001)	97.650 (0.339)	4.021 (0.003)	0.024 (0.003)
	SCAD	0.029 (0.002)	86.950 (0.753)	4.150 (0.011)	0.164 (0.011)
	MCP	0.032 (0.002)	86.850 (0.780)	4.176 (0.012)	0.177 (0.012)
	ENET	0.062 (0.003)	69.400 (1.031)	4.266 (0.018)	0.428 (0.017)
	LASSO	0.026 (0.002)	86.050 (0.775)	4.114 (0.010)	0.165 (0.010)
iv	Proposed	0.051 (0.003)	80.350 (0.889)	3.995 (0.005)	0.423 (0.021)
	SCAD	0.130 (0.004)	56.550 (1.109)	4.295 (0.018)	1.550 (0.039)
	MCP	0.195 (0.005)	43.500 (1.109)	4.191 (0.018)	1.550 (0.039)
	ENET	0.369 (0.005)	0.400 (0.141)	4.918 (0.062)	3.682 (0.037)
	LASSO	0.290 (0.004)	13.500 (0.7646)	4.957 (0.031)	2.200 (0.034)

Notation: FDR—false discovery rate; TRUE—percentage of the true model detected; SIZE—selected model size; HAM—Hamming distance; s.e.— standard error



**Fig. 1** Result of computational speed comparison: the proposed simultaneous marginal likelihood evaluation versus the for-loop

**Table 4** Model comparison results for BRCA data

	BIC	Extended BIC	MSPE
Proposed	985.93	1120.87	0.60
SCAD	1104.69	1176.42	0.68
MCP	1104.69	1176.42	0.68
ENET	1110.65	1198.68	0.68
LASSO	1104.69	1176.42	0.68

## 6 Real data application


In this section, we apply the proposed method to Breast invasive carcinoma (BRCA) data, which are generated by The Cancer Genome Atlas (TCGA) Research Network: <http://cancergenome.nih.gov>. We download BRCA data using R package `curatedTCGAData`. The data set contains 17,814 gene expression measurements (recorded on the log scale) of 526 patients with primary solid tumor in TCGA project. BRCA1 is a well-known tumor suppressor gene and its mutations predispose women to breast cancer (Findlay et al. 2018). Our goal here is to identify the best fitting model for estimating an association between BRCA1 (response variable) and the other genes (independent variables). After removing missing values, the data set reduces to  $n = 526$  samples with  $p = 17,326$  genes.

The proposed method and the penalized likelihood methods as in Sect. 5.1 are applied to the entire data of  $n = 526$  and  $p = 17,326$ . To assess model fitting, for each method, we compute BIC (Schwarz 1978), extended BIC, and mean squared prediction error (MSPE) for the ordinary least squares (OLS) estimate that is obtained by using the selected predictors, where MSPE is estimated by Monte Carlo cross-validation over 500 replications based on 70% of training set and 30% of testing set.

Table 4 summarizes model comparison results. As both BIC and extended BIC are minimized at the resulting model from the proposed method, this implies that our Bayesian method is most strongly supported by the data. In addition, the proposed method has the smallest MSPE. Hence, the results demonstrate that the proposed method provides the best fit to the data.

The heatmap in Fig. 2 displays the OLS coefficient estimates and the  $p$  values of the selected predictors for each method. As a result, the proposed method identifies 8 genes (GINS1, GNL1, VPS25, ARHGAP19, CRBN, NBR2, DTL, and KLHL13) that are statistically related to the human tumor suppressor gene, BRCA1.

According to the human gene database, called GeneCards (<https://www.genecards.org>), GINS1, CRBN, NBR2, DTL, and KLHL13 are known to be associated with diseases including neutropenia, cognitive disability, breast-ovarian cancer, myasthenic syndrome, and Zellweger syndrome. Krishnan et al. (2018) report that GNL1 is related to colorectal and gastric cancer. The study of Herz et al. (2006) shows that VPS25 can serve as a model for human cancer. In Marceaux et al. (2018), it is shown that the premature localization of ARHGAP19 at the cell membrane leads to cytokinesis failure and cell multinucleation.



**Fig. 2** A heatmap of OLS coefficient estimates for the selected predictors with  $p$  values without screening

## 7 Discussion

The proposed hybrid search approach is computationally beneficial and practically effective compared with traditional stochastic search algorithms that are commonly used in Bayesian variable selection. Although the idea of searching over the neighborhood is somewhat similar to that of the MCMC model composition ( $MC^3$ ) algorithm (Madigan and York 1995) and the shotgun stochastic search (SSS) algorithm (Hans et al. 2007), our search strategy is completely different from the existing methods. While  $MC^3$  and SSS perform a stochastic search over the model space of all possible candidates, we implement a hybrid search over the sub-model space of size  $k$ . In addition, evaluating multiple candidates for the next move can be done simultaneously in a single computation in the proposed method.

In addition, as mentioned in Sect. 4, parallel computing, which executes many calculations simultaneously using multiple nodes or multiple processors, can be employed in the proposed framework. The proposed hybrid search algorithm can be applied to frequentist model selection approaches such as Akaike information criterion (AIC), BIC, and Mallows'  $C_p$ . Our hybrid search algorithm can be also extended to multivariate regression and binary regression in a Bayesian framework. Such works are in progress by the authors.

**Acknowledgements** The authors are grateful to two anonymous reviewers for their helpful comments and constructive suggestions on an earlier version of this paper.

## Appendix

### A Calculation of Equation (3)

For any  $i \notin \hat{\mathcal{Y}}$ , from Eq. (2), we have

$$\begin{aligned} m(\mathbf{y}|\hat{\mathcal{Y}} \cup \{i\}) &\propto |\mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}}^{\top} \mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}} + \tau^{-1} \mathbf{I}_{k+1}|^{-1/2} \\ &\quad \times \left( \mathbf{y}^{\top} \mathbf{H}_{\hat{\mathcal{Y}} \cup \{i\}} \mathbf{y} + b_{\sigma} \right)^{-\frac{a_{\sigma} + n}{2}}. \end{aligned} \quad (13)$$

It follows from the Sherman–Morrison formula that

$$\mathbf{H}_{\hat{\mathcal{Y}} \cup \{i\}} = \mathbf{H}_{\hat{\mathcal{Y}}} - \frac{\mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{x}_i \mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}}}{\tau^{-1} + \mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{x}_i}. \quad (14)$$

Using the Sylvester’s determinant identity and the Sherman–Morrison formula, we obtain

$$\begin{aligned} &|\mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}}^{\top} \mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}} + \tau^{-1} \mathbf{I}_{k+1}| \\ &= \tau^{-(k+1)} |\tau \mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}} \mathbf{X}_{\hat{\mathcal{Y}} \cup \{i\}}^{\top} + \mathbf{I}_n| \\ &= \tau^{-(k+1)} |\tau \mathbf{X}_{\hat{\mathcal{Y}}} \mathbf{X}_{\hat{\mathcal{Y}}}^{\top} + \mathbf{I}_n + \tau \mathbf{x}_i \mathbf{x}_i^{\top}| \\ &= \tau^{-(k+1)} |\tau \mathbf{X}_{\hat{\mathcal{Y}}} \mathbf{X}_{\hat{\mathcal{Y}}}^{\top} + \mathbf{I}_n| \{1 + \tau \mathbf{x}_i^{\top} (\tau \mathbf{X}_{\hat{\mathcal{Y}}} \mathbf{X}_{\hat{\mathcal{Y}}}^{\top} + \mathbf{I}_n)^{-1} \mathbf{x}_i\} \\ &= |\mathbf{X}_{\hat{\mathcal{Y}}}^{\top} \mathbf{X}_{\hat{\mathcal{Y}}} + \tau^{-1} \mathbf{I}_k| (\tau^{-1} + \mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{x}_i). \end{aligned} \quad (15)$$

Applying (14) and (15) to (13), we thus have

$$\begin{aligned} m(\mathbf{y}|\hat{\mathcal{Y}} \cup \{i\}) &\propto \left\{ \mathbf{y}^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{y} - \frac{(\mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{y})^2}{\tau^{-1} + \mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{x}_i} + b_{\sigma} \right\}^{-\frac{a_{\sigma} + n}{2}} \\ &\quad \times (\tau^{-1} + \mathbf{x}_i^{\top} \mathbf{H}_{\hat{\mathcal{Y}}} \mathbf{x}_i)^{-1/2} \end{aligned}$$

for any  $i \notin \hat{\mathcal{Y}}$ .

### B Calculation of Eq. (4)

For any  $j \in \tilde{\mathcal{Y}}$ , Eq. (2) leads to

$$m(\mathbf{y}|\tilde{\mathcal{Y}} \setminus \{j\}) \propto |\mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}}^{\top} \mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}} + \tau^{-1} \mathbf{I}_k|^{-1/2}$$

$$\times \left( \mathbf{y}^\top \mathbf{H}_{\tilde{\mathcal{Y}} \setminus \{j\}} \mathbf{y} + b_\sigma \right)^{-\frac{a_\sigma + n}{2}}. \quad (16)$$

From the Sherman-Morrison formula, we have

$$\mathbf{H}_{\tilde{\mathcal{Y}} \setminus \{j\}} = \mathbf{H}_{\tilde{\mathcal{Y}}} + \frac{\mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{x}_j \mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}}}{\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{x}_j}. \quad (17)$$

From the Sylvester's determinant identity and the Sherman-Morrison formula, we obtain

$$\begin{aligned} & |\mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}}^\top \mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}} + \tau^{-1} \mathbf{I}_k| \\ &= \tau^{-k} |\tau \mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}} \mathbf{X}_{\tilde{\mathcal{Y}} \setminus \{j\}}^\top + \mathbf{I}_n| \\ &= \tau^{-k} |\tau \mathbf{X}_{\tilde{\mathcal{Y}}} \mathbf{X}_{\tilde{\mathcal{Y}}}^\top + \mathbf{I}_n - \tau \mathbf{x}_j \mathbf{x}_j^\top| \\ &= \tau^{-k} |\tau \mathbf{X}_{\tilde{\mathcal{Y}}} \mathbf{X}_{\tilde{\mathcal{Y}}}^\top + \mathbf{I}_n| \{1 - \tau \mathbf{x}_j^\top (\tau \mathbf{X}_{\tilde{\mathcal{Y}}} \mathbf{X}_{\tilde{\mathcal{Y}}}^\top + \mathbf{I}_n)^{-1} \mathbf{x}_j\} \\ &= \tau^2 |\mathbf{X}_{\tilde{\mathcal{Y}}}^\top \mathbf{X}_{\tilde{\mathcal{Y}}} + \tau^{-1} \mathbf{I}_{k+1}| (\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{x}_j). \end{aligned} \quad (18)$$

Hence, applying (17) and (18) to (16), we have

$$\begin{aligned} m(\mathbf{y} | \tilde{\mathcal{Y}} \setminus \{j\}) &\propto \left\{ \mathbf{y}^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{y} + \frac{(\mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{y})^2}{\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{x}_j} + b_\sigma \right\}^{-\frac{a_\sigma + n}{2}} \\ &\times (\tau^{-1} - \mathbf{x}_j^\top \mathbf{H}_{\tilde{\mathcal{Y}}} \mathbf{x}_j)^{-1/2} \end{aligned}$$

for any  $j \in \tilde{\mathcal{Y}}$ .

## References

- Bertsimas D, King A, Mazumder R (2016) Best subset selection via a modern optimization lens. *Ann Stat* 44(2):813–852
- Chen J, Chen Z (2008) Extended Bayesian information criteria for model selection with large model spaces. *Biometrika* 95(3):759–771
- Fan J, Li R (2001) Variable selection via nonconcave penalized likelihood and its oracle properties. *J Am Stat Assoc* 96(456):1348–1360
- Findlay GM, Daza RM, Martin B, Zhang MD, Leith AP, Gasperini M, Janizek JD, Huang X, Starita LM, Shendure J (2018) Accurate classification of BRCA1 variants with saturation genome editing. *Nature* 562(7726):217–222
- George EI, McCulloch RE (1993) Variable selection via Gibbs sampling. *J Am Stat Assoc* 88(423):881–889
- Hans C, Dobra A, West M (2007) Shotgun stochastic search for “large p” regression. *J Am Stat Assoc* 102(478):507–516
- Herz HM, Chen Z, Scherr H, Lackey M, Bolduc C, Bergmann A (2006) vps25 mosaics display non-autonomous cell survival and overgrowth, and autonomous apoptosis. *Development* 133(10):1871–1880
- Hocking RR, Leslie RN (1967) Selection of the best subset in regression analysis. *Technometrics* 9(4):531–540
- Kass RE, Raftery AE (1995) Bayes factors. *J Am Stat Assoc* 90(430):773–795



- Kirkpatrick S, Gelatt CD, Vecchi MP (1983) Optimization by simulated annealing. *Science* 220(4598):671–680
- Krishnan R, Boddapati N, Mahalingam S (2018) Interplay between human nucleolar gnl1 and rps20 is critical to modulate cell proliferation. *Sci Rep* 8:11421
- Liang F, Liu C, Carroll RJ (2007) Stochastic approximation in Monte Carlo computation. *J Am Stat Assoc* 102(477):305–320
- Liang F, Song Q, Yu K (2013) Bayesian subset modeling for high-dimensional generalized linear models. *J Am Stat Assoc* 108(502):589–606
- Madigan D, York J (1995) Bayesian graphical models for discrete data. *Int Stat Rev* 63(2):215–232
- Marceaux C, Petit D, Bertoglio J, David MD (2018) Phosphorylation of arhgap19 by cdk1 and rock regulates its subcellular localization and function during mitosis. *J Cell Sci.* <https://doi.org/10.1242/jcs.208397>
- Schwarz G (1978) Estimating the dimension of a model. *Ann Stat* 6(2):461–464
- Tibshirani R (1996) Regression shrinkage and selection via the lasso. *J R Stat Soc Ser B (Stat Methodol)* 58(1):267–288
- Zhang CH (2010) Nearly unbiased variable selection under minimax concave penalty. *Ann Stat* 38(2):894–942
- Zou H, Hastie T (2005) Regularization and variable selection via the elastic net. *J R Stat Soc Ser B (Stat Methodol)* 67(2):301–320

**Publisher's Note** Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.